FutureSight

Technical Takeaway Assignment: ReviewLens AI

Business Context

Imagine a consultancy that specializes in Online Reputation Management (ORM). Their business model relies on analyzing massive amounts of fragmented customer feedback to offer strategic services to brands.

Currently, their analysts spend hours manually reading reviews to identify "pain points." They need a rapid prototype of a Review Intelligence Portal that can ingest a product's digital footprint and allow an analyst to "talk" to that data to find specific trends—without the AI drifting into generalities or competitor data.

The Mission

Develop ReviewLens AI: A secure, web-based portal that enables a user to track a product or entity from a single review platform (Amazon, Google Maps, G2, Capterra, or a similar publicly accessible platform) and analyze those reviews using a guardrailed Q&A interface.

Whichever platform you choose, it should:

- Be publicly accessible — open to browse without authentication.
- Feature user-generated content — customer-written text reviews and ratings.

How We Evaluate

We aren't testing whether you can code without AI — we're testing whether you can build great software with it. Using AI tools (Claude Code, Cursor, Copilot, etc.) is expected and strongly encouraged.

As a senior candidate, we're evaluating your ability to ship a professional product quickly, your judgment in how you work with AI, and the engineering instincts you bring to the process.

Core Requirements

1. Ingestion & Scraping Summary

- Ingestion Module: The application should accept a target URL from the chosen platform and extract the relevant review data. If scraping proves impractical (anti-bot measures, dynamic rendering, etc.), an alternative import mechanism (e.g., CSV upload, paste) is a perfectly valid approach.
- Scraping Summary: Give the user a clear picture of what was captured — whether that's text-based, a table, or a visual dashboard. The goal is confidence that the data is accurate, complete, and ready for analysis.

2. Guardrailed Q&A Interface

- Interactive Chat: Build an interface where users can pose questions exclusively about the ingested reviews.
- Scope Guard Enforcement: This is one we care a lot about. If a user asks about an external platform or general world knowledge, the AI should gracefully and explicitly decline (e.g., if

tracking Google Maps, it shouldn't discuss Amazon reviews or the current weather). This should be primarily driven by your system prompt configuration.

3. Deployment

- Hosting: We'd like to see the application hosted publicly and accessible via a URL.
- Code: Please share the full source code in a GitHub Repository.

Make This Your Own

This assignment is your canvas. The core requirements above set the baseline, but this is where you can go above and beyond to stand out. Take this project in whatever direction proves your unique value — whether that's through sophisticated prompt engineering, an elegant UI, a particularly clever architectural choice, or engineering practices that show how you'd operate on a production team. Surprise us.

Constraints & Out-of-Scope

- Zero Financial Cost: We ask that this be completed at zero cost — no reimbursements will be provided. Free-tier services are fair game.
- Tooling: You're welcome to use new email addresses for "Free Trial" tiers of tools/LLMs to maintain zero cost. This is a possible path, not a mandate, as long as you can deliver the project under the cost constraint.
- No User Auth: The application should be directly accessible via its URL without a login.

Deliverables

1. GitHub Repository: A link to the complete code, shared with the hiring team.
2. **AI Session Transcripts: Include your full AI session transcripts (Claude Code session logs, Cursor Composer history, Copilot chat exports, etc.) in the repo under an /ai-transcripts directory. This is a first-class deliverable. Don't clean them up or cherry-pick — we want the real working session, including dead ends.**
3. Live URL: A direct link to your deployed application.
4. Loom Demo (under 3 minutes): Walk us through what you built:
    - Demo the full flow end-to-end (scraping → summary → Q&A → scope guard).
    - Talk through your key design decisions and tradeoffs.
    - Call out anything you're proud of, and anything you'd do differently with more time.
5. README.md: Setup instructions, architecture overview, and any assumptions you made.

Final Instructions

- Assumptions: Use your best professional judgment for any detail not explicitly covered. Anything not mentioned should be assumed and proceeded with.
- No Further Guidelines: To simulate a real-world rapid deployment, no additional clarifications will be provided after the assignment is issued.

- Time: We suggest 5 hours as a planning target. This is not a strict limit — if you want to invest more time to go deeper, that's your call.